

CS704: Lecture 3

Lambda Calculus: Substitution, Reduction, and Normal Forms

Thomas Reps

[Based on notes taken by Tycho Andersen on January 25, 2010]

Abstract

This lecture concerns the rules for computing on terms in the λ -calculus. It discusses substitution, reduction rules, and normal forms.

1 Introduction and Review

The λ -calculus model of computation was invented by Alonzo Church in the 1930's. In today's terms, it can be thought of as a (very!) stripped-down functional programming language. It is a *declarative* (as opposed to *imperative*) programming language: there is no notion of assignment. In fact, all computation is performed by rewriting one λ -term to another. Roughly, a computation halts (it if does) when no more rewriting is possible. Despite the stripped-down nature of the language, λ -calculus is Turing-complete. That is, any problem that can be formulated as a Turing-machine computation can be formulated as a λ -calculus computation, and vice versa.

Last time we gave a deliberately imprecise rewriting rule (the “naive” rule of β -reduction):

$(\lambda x.M)N \rightarrow_{\beta} M'$, where M' is M with all occurrences of x in M replaced by N .

In English, given a function applied to an argument, rewrite it to the function body (which is $x + 1$ in the example above) with all occurrences of the formal (x above) replaced with the argument (3 above).

2 How Does λ -Calculus Stack Up as a Computational System

To see how the λ -calculus stacks up as a computational system, let's ask (and answer) a few rhetorical questions.


Does every λ -term have a normal form? No, there exist λ -terms that do not have a normal form. In particular, consider the following two examples:

$$(\lambda x.xx)(\lambda x.xx) \rightarrow_{\beta} (\lambda x.xx)(\lambda x.xx) \rightarrow_{\beta} \dots$$

$$(\lambda x.xxx)(\lambda x.xxx) \rightarrow_{\beta} (\lambda x.xxx)(\lambda x.xxx)(\lambda x.xxx) \rightarrow_{\beta} \dots$$

In the first example, at each step the term rewrites to itself, creating an infinite loop. In the second example, an extra copy of $(\lambda x.xxx)$ is introduced at each step, so the term grows in size with each “reduction” step. (Thus, reduction is a misnomer, because the *size* of the λ -term that results from a “reduction” step does not necessarily decrease.)


If a λ -term has a normal form, does every choice of reductions lead to that normal form? No, consider the following λ -terms $W \stackrel{\text{def}}{=} (\lambda x.xxx)(\lambda x.xxx)$, $F \stackrel{\text{def}}{=} \lambda x.\lambda y.y$, $I \stackrel{\text{def}}{=} \lambda x.x$, $J \stackrel{\text{def}}{=} \lambda y.y$, and the transitions possible from the λ -term $F(WW)I$.



If one always chooses to perform a reduction step that rewrites an occurrence of WW , the reduction sequence never terminates. However, at each stage each λ -term in the chain on the left-hand side can, in two reduction steps, be reduced to a normal form, as shown in the diagram.

The example relates to the programming-language notion of *strictness*: a function is *strict* in a given argument position i if whenever the i^{th} actual parameter diverges, the function diverges. An abstraction in λ -calculus, therefore, is *not* strict in its (one) argument.

If a λ -term has a normal form, is there a reduction strategy that will always find it? Yes, by repeatedly performing the leftmost-outermost reduction. The leftmost-outermost β -redex can be found by simply doing a pre-order traversal of the abstract-syntax tree until the first β -redex is encountered. Below is a diagram that shows a λ -term that has a total of four redexes, two of which are outermost redexes. The one circled with a solid line is the leftmost-outermost redex.



Applied λ -calculi. Other languages can be considered as “sugared” (or “applied”) λ -calculi [1]. In “pure” λ -calculus, there are no primitive data values or functions, but applied λ -calculi can be created by adding

- constant symbols (e.g., 0, 1, 2, ...; 'a', 'b', ...; "", "a", "aa", ...)
- constant functions (e.g., "+", "*", ...)

(as well as possibly other elements).

An example of a term in an applied λ -calculus is $\lambda x.x + 1$. Again, to compute one would repeatedly find occurrences of an application of a λ -term (say, $\lambda x.x + a$) to an argument (which is itself also a λ -term). Computation is performed by repeatedly performing a rewriting step on the current λ -term until one is left with a term that cannot be rewritten any further, e.g.,

$$(\lambda x.x + 1)3 = 3 + 1 = 4$$

Is the fact that λ -calculus has only one-argument functions a limitation? No, multi-argument functions can be built up by using nested abstractions. For instance, $\lambda x.\lambda y.x$ is an example of a “two-argument” function built by nesting $\lambda y.x$ inside of $(\lambda x.\dots)$. $\lambda x.\lambda y.x$ is a sort of “first” function. That is, it takes two actual parameters and returns the first of them, ignoring the value of the second actual parameter.

λ -calculus allows familiar functions to arise in unfamiliar ways. In an applied λ -calculus, the “add” function is $\lambda x.\lambda y.x + y$. This “two-argument” function can be used to form more specific functions, such as the “add three” function, by rewriting:

$$(\lambda x.\lambda y.x + y)3 = \lambda y.3 + y$$

The example illustrates that in λ -calculus the add function is a *higher-order function*: it takes an actual parameter and gives back a function as the result. (Another kind of higher-order function takes a function as an actual parameter. In λ -calculus, all functions are higher-order in this sense.)

Rewriting in the λ -calculus provides interesting ways of manipulating programs. Suppose that instead of $\lambda y.3 + y$, we wanted to obtain $\lambda y.y + 4$ from the add function $\lambda x.\lambda y.x + y$. How can we obtain this? Note that

$$(\lambda x.\lambda y.x + y)4 = \lambda y.4 + y,$$

which is different from the desired result.

We can create the desired term by using a *combinator* (i.e., a λ -term with no free (unbound) variables). If we apply the combinator $\lambda f.\lambda a.\lambda b.fba$ to the add function, and then apply the resulting λ -term to 4 we have:

$$\begin{aligned} ((\lambda f.\lambda a.\lambda b.fba)(\lambda x.\lambda y.x + y))4 &\rightarrow (\lambda a.\lambda b.(\lambda x.\lambda y.x + y)ba)4 \\ &\rightarrow (\lambda a.\lambda b.(\lambda y.b + y)a)4 \\ &\rightarrow (\lambda a.\lambda b.b + a)4 \\ &\rightarrow \lambda b.b + 4 \end{aligned}$$

Remark. Combinators will play a role in the next lecture when we introduce *encodings* that (i) simulate data of various kinds, (ii) operations on that data, and (iii) various programming constructs (such as conditional expressions and recursion). \square

Reasoning about programs. Rewriting (reduction, substitution of equals for equals) gives us a tool for reasoning about programs. Such reasoning is difficult for other languages if either (i) it has notions of assignment, destructive updating, etc., or (ii) all you have is a “definition” of the language based on what is produced by a compiler.

Recursion. The λ -term $\lambda x.x$ represents the identity function, and hence is roughly equivalent to the (unnamed) C function defined by

```
datum ?(datum x) {
    return x;
}
```

What corresponds to

```
int fact(int x) {
    return(x == 0 ? 1 : x * fact(x-1));
}
```

The answer is that there is an analog of recursive definitions in the λ -calculus, but in pure λ -calculus the problem is that functional expressions do not have permanent names (as **fact** does above). That feature is key for recursion to work in C: in the body of **fact**, the name **fact** is used to make the recursive call.

In λ -calculus, a combinator will let us solve the problem of not having permanent names. (In essence, the combinator’s bound variables essentially “issue” temporary names.)

3 A Precise Definition of β -Reduction

First, some terminology:

- Scoping: With which enclosing occurrence of a λ is each variable associated? (In λ -calculus, one rarely talks of scopes, but in conventional programming-language terms, each occurrence of $(\lambda x.\dots)$ creates a scope and the λx part is like a declaration of x that is added for that scope.)
- Bound variable: A variable that is associated with an enclosing λ .
- Free variable: A variable that is *not* associated with an enclosing λ .

Note that it is important to be clear about which (sub-)term is being referred to. For instance, consider a λ -term of the form $P \stackrel{\text{def}}{=} \lambda x.M$, and assume that none of the abstractions inside of M are of the form $(\lambda x.\dots)$. The occurrences of variable x in P are *bound variables in P* . However, if we are just considering M alone, the occurrences of x in M are *free variables in M* .

Definition 3.1 (Free and bound variables) *The free and bound variables of a λ -term are defined in terms of the structure of the term, as follows:*

1. In an expression of the form x (i.e., a variable),

- x is free
 - there are no bound variables
2. In an expression of the form $\lambda x.M$:
- every x in M is bound
 - for every other variable y , if y is free in M , then it is free in $\lambda x.M$; if y is bound in M , then it is bound in $\lambda x.M$.
- (Note that a variable y can be both free and bound in M : some occurrences of y can be free in M , while if M has a sub-term of the form $(\lambda y.\dots)$, other occurrences of y can be bound in M .)
3. In an expression of the form MN : this is the same as in 2, but for both M and N .

Alternatively, the free variables can be defined inductively as follows:

$$\begin{aligned} FV(x) &\stackrel{\text{def}}{=} \{x\} \\ FV(\lambda x.M) &\stackrel{\text{def}}{=} FV(M) - \{x\} \\ FV(M N) &\stackrel{\text{def}}{=} FV(M) \cup FV(N) \end{aligned}$$

□

Next, we introduce the correct β -reduction rule. The naive β -reduction rule has some problems because it does not respect scoping.

- Sometimes we do not want to replace *all* the occurrences of the formal in the function body. In the example below, the variable x is overloaded, and replacing all occurrences yields the wrong answer:

$$\begin{aligned} (\lambda x.x + ((\lambda x.x + 1)(3)))2 &= 2 + ((\lambda x.2 + 1)(3)) && \text{(Bad!)} \\ &= 2 + 2 + 1 \\ &= 5 \end{aligned}$$

To avoid this problem, we will use a rewriting rule called α -reduction, which will be used to rename formal parameters. Below is the correct evaluation.

$$\begin{aligned} (\lambda x.x + ((\lambda x.x + 1)(3)))2 &= 2 + ((\lambda z.z + 1)(3)) && \text{(Ok!)} \\ &= 2 + (3 + 1) \\ &= 6 \end{aligned}$$

- The naive β -reduction rule permits a variable that is free in an argument to become bound. Consider $\lambda x.\lambda y.x$, which is a two-argument function that returns the first argument and ignores the second. If the variable names are not overloaded (i.e., the actual parameter(s) do not have the same name(s) as a formal parameter(s) in a nested function(s)), we can successfully evaluate this function using the naive rule:

$$((\lambda x.\lambda y.x)w)z = (\lambda y.w)z = w$$

However, if variable names are overloaded, we may obtain the wrong answer.

$$((\lambda x.\lambda y.x)y)z = (\lambda y.y)z = z$$

α -reduction is applied to abstractions. In $\lambda x.M$ it replaces all *free* occurrences of x in M with some other identifier that is not in M . Some examples:

$$\begin{aligned}\lambda x.(\lambda y.x) &\rightarrow_\alpha \lambda z.\lambda y.\lambda z \\ \lambda x.(\lambda x.x) &\rightarrow_\alpha \lambda z.\lambda x.\lambda x\end{aligned}$$

Formally, the α -reduction and β -reduction rules are defined in terms of the operation of *substitution* defined in Defn. 3.2.

Remark. I use several kinds of equality-like symbols:

- $P \stackrel{\text{def}}{=} Q$ means that P is defined to be the λ -term Q .
- $P \equiv Q$ means that P and Q are identical. Thus, $P \not\equiv Q$ means that there is some difference between P and Q .

□

Definition 3.2 (Substitution) *The substitution of a term N for all free occurrences of a variable x in M , denoted by $M[N/x]$,¹ is defined inductively as follows, according to the structure of M :*

1. $x[N/x] \stackrel{\text{def}}{=} N$
2. $y[N/x] \stackrel{\text{def}}{=} y$ iff $y \neq x$
3. $MP[N/x] \stackrel{\text{def}}{=} (M[N/x]P[N/x])$
4. $(\lambda x.M)[N/x] \stackrel{\text{def}}{=} \lambda x.M$
5. $(\lambda y.M)[N/x] \stackrel{\text{def}}{=} \lambda z.((M[z/y])[N/x])$, where $y \neq x$ and z is fresh

□

Definition 3.3 (Reduction Rules) *The rules of α -reduction and β -reduction are defined as follows:*

α -reduction: $\lambda x.M \rightarrow_\alpha \lambda y.M[y/x]$
 β -reduction: $(\lambda x.M)N \rightarrow_\beta M[N/x]$.

□

References

- [1] P. Landin. The next 700 programming languages. *Commun. ACM*, 9(3):157–166, 1966.

¹ $M[N/x]$ is properly read as “ M with N substituted for all free occurrences of x ”; however, it is often shortened to “ M with N in for x ”.